

Lecture 1: Gale-Shapley Algorithm for Stable Matching

Stable Matching Problem: There are two groups of people: $A = \{a_1 \dots a_n\}$ and $B = \{b_1 \dots b_n\}$, and we want to find a stable matching between A and B such that there is no pair of people a_i, b_j who would rather be with each other than their current partners. i.e We want to find a permutation π such that there are no $i, j \in [n]$ for whom

1. $\pi(i) \neq j$
2. a_i prefers b_j to $b_{\pi(i)}$
3. b_j prefers a_i to $b_{\pi^{-1}(j)}$

Gale-Shapley algorithm: The people in group A make offers in the order of their preference lists. When b_j receives an offer from a_i , if b_j is unmatched or prefers a_i to their current partner, they accept and become partners with a_i . Otherwise, b_j rejects and a_i moves onto the next person.

Properties:

- The Gale-Shapley algorithm terminates in $O(n^2)$ steps and results in a stable matching.

Proof: There are n^2 possible offers and each offer is made at most once.

Terminates: Observe that if some a_i makes an offer to their last choice b_j , then all people except b_j must be unavailable for a_i and thus already matched. There are only $n - 1$ other people in A , so b_j must be unmatched, and b_j accepts a_i 's offer, and the algorithm terminates.

Stable: Assume there are $i, j \in [n]$ such that a_i and b_j prefer each other to their current partners. Let b_j be a_i 's current partner and b_j' be matched with a_i . Then b_j is available to a_i . However, since a_i is matched with b_j' , b_j must have been unavailable to a_i at some point, a contradiction.

- As Gale-Shapley progresses, (1) people in B only become happier, and (2) if b_j becomes unavailable to a_i they never become available again.

Proof: (1) b_j only breaks off from its match if it prefers the new a_i . (2) Assume b_j becomes unavailable to a_i and then becomes available again. Let a_i' be b_j 's partner when b_j became unavailable to a_i . Let a_i'' be b_j 's partner when b_j became available to a_i again. Then, b_j prefers a_i' to a_i and a_i to a_i'' . However, b_j must prefer a_i'' to a_i' , a contradiction.

Lecture 2: Big O Notation

Big O Notation: Given functions $f, g : \mathbb{N} \mapsto \mathbb{R}^+$

1. We say that $f(n)$ is $O(g(n))$ if there exist $n_0, C > 0$ such that for all $n \geq n_0$, $f(n) \leq Cg(n)$
2. We say that $f(n)$ is $\Omega(g(n))$ if there exist $n_0, c > 0$ such that for all $n \geq n_0$, $f(n) \geq cg(n)$
3. We say that $f(n)$ is $\Theta(g(n))$ if $f(n)$ is $O(g(n))$ and $\Omega(g(n))$ such that there exist n_0, c, C' such that for all $n \geq n_0$, $cg(n) \leq f(n) \leq C'g(n)$

Examples:

- If $f(n) \leq 8n \log_2(n) + 20n + 100$ then $f(n)$ is $O(n \log n)$
- If $f(n) \geq n^2 - 3n - 2$ then $f(n)$ is $\Omega(n^2)$
- If $\frac{1}{2} \log_2 n - 2 \leq f(n) \leq 4 \log_2 n + 1$ then $f(n)$ is $\Theta(\log n)$

Properties:

- Given $f, g, h : \mathbb{N} \mapsto \mathbb{R}^+$, if $f(n)$ is $O(g(n))$ and $g(n)$ is $O(h(n))$, then $f(n)$ is $O(h(n))$
- Given $f_1, f_2, g_1, g_2 : \mathbb{N} \mapsto \mathbb{R}^+$, if $f_1(n)$ is $O(g_1(n))$ and $f_2(n)$ is $O(g_2(n))$ then:

- $f_1(n) + f_2(n)$ is $O(g_1(n) + g_2(n))$ and $O(\max\{g_1(n), g_2(n)\})$
- $f_1(n)f_2(n)$ is $O(g_1(n)g_2(n))$

Proof: There exist n_1, C_1 such that $\forall n \geq n_1 (f_1(n) \leq C_1 g_1(n))$ and there exist n_2, C_2 such that $\forall n \geq n_2 (f_2(n) \leq C_2 g_2(n))$. Let $n' = \max\{n_1, n_2\}$ and $C' = C_1 C_2$, So for all $n \geq n'$, $f_1(n)f_2(n) \leq C' g_1(n)g_2(n)$

Comparing logarithms, polynomials, and exponential functions:

1. $\log n^{C'}$ is $O(n^C)$
2. $n^{O(1)}$ means at most n^C for some $C > 0$. This is polynomial time
3. $2^{O(n)}$ means at most 2^{Cn} for some $C > 0$. This is exponential time.

Lecture 3: Greedy Algorithms

Interval Scheduling Problem 1 (Maximizing number of jobs): Given n jobs with a set time interval $[a_i, b_i]$ and one processor, find a schedule S which accepts the maximum number of jobs without having two jobs running at the same time. Formally, we say that a sequence $S = (i_1, \dots, i_m)$ is a valid schedule if for all $j \in [m-1]$ ($b_{i_j} \leq a_{i_{j+1}}$)

Greedy Algorithm for Interval Scheduling 1: When choosing the next job, always choose the job which can be finished first.

- *Stored Data*: $S_k = (i_1, \dots, i_k)$ of the jobs which have been accepted thus far, and the time $t_k = b_{i_k}$
- *Initialization*: $S_0 = \emptyset$ and $t_0 = 0$
- *Iterative Step*: Choose the next job i_{k+1} from the set $\{i \in [n] \mid a_i \geq t_k = b_{i_k}\}$ such that b_{i_k} is minimized. Update S_{k+1} and $t_{k+1} = b_{i_{k+1}}$, if there are no available jobs left then the algorithm terminates and we take $S = S_k$

Properties:

- Given a valid schedule $S = (i_1, \dots, i_m)$, for all $j \in [m]$, we denote the time at which the j th job in S finishes as $t_j(S) = b_{i_j}$. For all $j > m$, we say $t_j(S) = \infty$
- If the greedy algorithm gives a valid schedule $S = (i_1, \dots, i_m)$, then for any other valid schedule $S' = (i_1, \dots, i'_m)$, $m' \leq m$
Proof (Using the lemma proved below): Consider $j = m+1$, $t_{m+1}(S) = \infty$ so we must have that $t_{m+1}(S') = \infty$. Thus, $m' \leq m$
- If $S = (i_1, \dots, i_m)$ is the schedule given by the greedy algorithm then for any other valid schedule $S' = (i_1, \dots, i'_m)$, for all $j \in \mathbb{N}$, either $t_j(S) \leq t_j(S')$ or $t_j(S') = \infty$
Proof: By induction on j . For the inductive step, assume that $t_k(S) \leq t_k(S')$. Note that $a_{i'_{k+1}} \geq t_k(S') \geq t_k(S)$ so job i'_{k+1} is available to S . Since S always chooses the job with the earliest completion time out of the available jobs, $t_{k+1}(S) = b_{i_{k+1}} \leq b_{i'_{k+1}} = t_{k+1}(S')$.
- This algorithm can be implemented in $O(n \log n)$ time by first sorting the jobs in increasing order of b_i

Interval Scheduling Problem 2 (Minimizing maximum lateness): Given n jobs, each of which has length $l_i > 0$ and a deadline d_i , complete the jobs one by one such that the maximum lateness of any single job is minimized. Formally, define $t_S(i_k) = \sum_{j=1}^k l_{i_j}$ to be the time at which job i_k is finished, and find a schedule $S = (i_1, \dots, i_n)$ such that $\max_{j \in [n]} \{t_S(i_j) - d_{i_j}\}$ is minimized.

Greedy Algorithm for Interval Scheduling 2 (Earliest Deadline First): Take the schedule S_{greedy} such that d_{i_1}, \dots, d_{i_n} are in increasing order.

Properties:

- The schedule S such that d_{i_1}, \dots, d_{i_n} are in increasing order minimizes $\max_{j \in [n]} \{t_S(i_j) - d_{i_j}\}$
Proof: Given an optimal schedule S that is not S_{greedy} , we can obtain S_{greedy} by exchanging pairs of jobs without increasing maximum lateness. Therefore, S_{greedy} is optimal.
- If S is an optimal schedule, and $d_{i_j} > d_{i_{j+1}}$ then the schedule S' obtained by swapping i_j and i_{j+1} is also optimal.
Proof:
 1. Swapping i_j and i_{j+1} does not affect the lateness of any other job as for all $j' \in [n] \setminus \{j, j+1\}$, $t_{S'}(i_{j'}) = t_S(i_{j'})$
 2. $t_S(i_{j+1}) = t_{S'}(i_j) = t_{S'}(i_{j+1}) + l_{i_j}$ as schedule S' puts job i_{j+1} before job i_j
This implies that $t_S(i_{j+1}) - d_{i_{j+1}} > t_{S'}(i_{j+1}) - d_{i_{j+1}}$ and $t_S(i_{j+1}) - d_{i_{j+1}} > t_{S'}(i_j) - d_{i_j}$ as $d_{i_j} > d_{i_{j+1}}$. Thus, the lateness of job i_{j+1} for schedule S is greater than the lateness of both job i_j and i_{j+1} for schedule S' . So the maximum lateness for S' is less than or equal to the maximum lateness of S .

Lecture 4: Dijkstra's Algorithm

Shortest Path Problem: Given a directed graph $G = (V, E)$ with edge lengths $\{\ell_{uv} : (u, v) \in E\}$, a starting vertex s , and a destination vertex t , find the shortest path from s , to t .

Properties:

- If the edge lengths are non-negative, at each step, Dijkstra's algorithm correctly computes the distance from s to each vertex $u \in S$
 Proof: By induction on $..$ Inductive step: assume the distances $\{d_u : u \in S\}$ are correct. Let (u, v) be the next edge that Dijkstra's algorithm considers where $v \notin S$. This gives a path P of length $d_u + \ell_{uv}$. If P' is another path from s to v , let (u', v') be the first edge on P' leaving s . $d_{u'} + \ell_{u'v'} \geq d_u + \ell_{uv}$, so the length of P' is greater than or equal to P .
-

Algorithm 1 Dijkstra's Algorithm

Lecture 6: Minimum Spanning Trees

Minimum Spanning Tree: Given a graph $G = (V, E)$, the minimum spanning tree is a set of edges $T \subseteq E$ such that (V, T) is a tree and $\sum_{e \in T} \ell_e$ is minimized.

Kruskal's Algorithm: Sort the edges E in order of their length. We then go through the edges one by one and take each edge which does not form a cycle.

- Stored Data: A set T of the edges taken thus far
- Iterative Step: We take the next edge $e \in E$. If $T \cup \{e\}$ contains a cycle, discard e , otherwise add e to T . The algorithm terminates when $|T| = n - 1$ or there are no edges left to consider

Determining if $T \cup \{e\}$ creates a cycle can be done in $O(\alpha(n))$ time via union-find where $\alpha(n)$ is the inverse Ackermann function.

Prim's Algorithm: Iteratively accepts the shortest edge which leads to a new vertex.

- Stored Data: A set T of edges and a set S of vertices taken thus far, and a set of edges $E_{\text{candidate}}$ with their lengths ℓ_e
- Iterative Step: Remove the shortest edge $e = \{v, w\} \in E_{\text{candidate}}$. If w is not in S , then we add w to S and e to T , and update $E_{\text{candidate}}$ with all of w 's edges. The algorithm terminates when $S = V$ or $E_{\text{candidate}}$ is empty (in which case G is not connected).

A cut is a partition (L, R) of the vertices of G where $L \neq \emptyset$ and $R \neq \emptyset$. We say an edge $e = \{u, v\}$ crosses the cut (L, R) if u is in L and v is in R or vice versa.

Theorem: For each edge $e \in E$, e is in the MST if and only if there exist a cut (L, R) such that e is the shortest edge crossing (L, R)

Proof: If e is the shortest edge crossing (L, R) , assume $e \notin T$. If so, adding e creates a cycle, which contains some other edge e' crossing (L, R) . So $T' = (T \cup \{e\}) \setminus \{e'\}$ has shorter total length than T , a contradiction.

If $e = \{u, v\} \in T$, take L to be the set of vertices reachable from u using the edges $T \setminus \{e\}$. If there were a shorter edge e' crossing (L, R) , $T' = (T \cup \{e'\}) \setminus \{e\}$ would have shorter total length.

Theorem: For all edges $e \in E$, $e \notin T$ if and only if there exists a cycle C such that e is the longest edge of C

Correctness of Kruskal's Algorithm: If e is the shortest edge crossing a (L, R) , adding e cannot create a cycle so e will be taken. If Kruskal's algorithm takes $e = \{u, v\}$, let T be the set of edges which were taken so far. Take L to be the set of vertices reachable from u with the edges in T . There cannot be a shorter edge e' crossing (L, R) , as otherwise e' would've been considered before e .

Correctness of Prim's Algorithm: Observe that at each step, Prim's algorithm takes the shortest edge crossing the cut between S and the remaining vertices. So Prim's algorithm only takes edges in T . If G is connected, the output of Prim's algorithm must be connected. So all edges of T must be taken.